

ALICE O² data model proposal

CWG4

March 2, 2017

Abstract

This note presents a proposal for the ALICE O² data model. A base data layout and metadata format that allow for efficient resource use are proposed. Application of the data model to online/offline data processing and quality control is discussed.

1 Introduction

The ALICE online-offline (O²) computing system [1, 2] is a computing facility and a software framework designed for the processing of the ALICE data in the upcoming LHC Run 3. The design aims at high data throughput and parallelism using a multiprocess model. It does not, however, exclude the use of multithreading and other forms of concurrent processing inside of individual processes.

The data exchange between processes running within the O² system (called O² devices) is taken care of by the ALICE-FAIR (Alfa) framework [3]. Since this is the main communication mechanism foreseen for data exchange, it effectively serves the role of an API between the devices. The Alfa framework provides data transport and synchronisation primitives via the FairMQ message queue library. FairMQ messages consist of raw memory buffers which are asynchronously queued and atomically delivered.

The online data processed by the O² system consists of data buffers originating from the detector hardware (raw data) and the processing devices (containing derived data). The data fragments belonging to a logical unit are grouped into a data set, e.g. a (sub-) time frames. A (sub-) time frame contains raw data associated to a period of data taking (typically several tens of ms, as dictated by the heartbeat trigger [4]) and/or the results of the processing of these data. In addition, any data that might be necessary to describe the data set can be added to the logical group.

In the unified online-offline software model also derived data is handled within the same software framework. O² devices dedicated to quality control (QC) and physics analysis tasks should use the same set of interfaces as components in the synchronous data reconstruction and calibration chain. The requirements for the data used in these tasks tend to be different from the online components: high level abstractions and ease of use (of e.g. ROOT[5] objects), despite the additional overhead, is sometimes preferred to high performance low-level data structures. Transparent support for high level data structures is part of the proposed data model.

2 Zero-copy approach

A single uncompressed time frame data volume is expected to be of the order of tens of gigabytes. The data model should facilitate data exchange approaches that minimize resource use, i.e. avoid unnecessary copies of data and serialization/de-serialization overhead.

Traditionally, in monolithic designs where all processing is performed inside a single process, data exchange between logical processing units (i.e. subroutines and threads) is supported at the programming language semantics level using parameter passing (by value or by reference). Subroutines can efficiently exchange data by passing references (pointers) to resources available in the uniform virtual memory space available to a process. A flat process memory model also enables fully transparent and efficient compiler support of higher level abstractions like virtual inheritance in the C++ language, usually implemented as (hidden) references to internal metadata (e.g. a virtual method table).

In the O^2 multi-processing model data needs to be exchanged between devices running as separate processes. This scenario is not covered by the core language (in the case of C++) and needs to be addressed using additional abstraction layers. In the most general case the data to be shared is copied into a transport buffer a copy of which is made available to the receiving process by the operating system using various interprocess or network communication mechanisms e.g. pipes or (network) sockets. If the data structure is complex, i.e. not contiguous and/or refers to other data structures (contains pointer/reference members) it needs to be copied (serialized) to a contiguous buffer including the metadata needed to recreate the structure in the receiving process memory. The buffer then has to be de-serialized by each consumer.

Various libraries exist that provide this functionality and try to optimize various aspects of the serialization and de-serialization process, e.g. google flatbuffers[6] allow zero-copy access to data in the serialized buffer with only a minimal (indirection) overhead to ensure portability and schema evolution. Serialization, however always incurs CPU overhead and data duplication as additional resources need to be allocated for the target buffer.

In the case of devices separated by a network interface (e.g. running on different machines in a network) a data copy over the network is inevitable. The copy overhead can, however, be minimized using modern networking technologies utilizing remote direct memory access (RDMA).

For processes running on the same machine it is possible to use shared memory which can be accessed directly by more than one process. Shared memory segments are mapped by the operating system at arbitrary base addresses within the process virtual memory. Data access in such a memory segment therefore cannot rely on absolute references as they depend on the specific process address space layout and are valid in one process only. That excludes straightforward use of structures that contain reference (pointer) data members and/or higher level language features (like virtual inheritance) as these in general are implemented using internal references. Shared memory regions are usually not covered by the standard resource allocation mechanisms provided by the language and memory management needs to be handled via additional libraries.

From the communication, or data transport perspective the use of either transport buffers for socket-like communication or special memory regions is similar in the sense that in both cases a contiguous data buffer

is received, or accessed, at a random base address, imposing similar limitations on the data representation. The choice is to either serialize/deserialize high level data for transport when the cost of this process is acceptable, or avoid data copies and the CPU overhead by constructing the data in a fully self-contained format directly in a buffer that is suitable for the selected transport mechanism. In the latter case, since the transport mechanism is chosen at run time, all memory allocations need to be performed via the transport layer explicitly.

Additional constraints on data layout and representation are given by the machine architecture. True zero-copy approach is only possible if all machines involved in the processing have the same internal binary data representation with regards to endianness and data structure alignment as the agreed format. In this case data produced on one machine can be accessed by an other directly from the transport buffer. Otherwise data needs to be converted to a suitable binary representation; this can in principle be handled transparently by either the transport layer or the data access interface.

In view of the large volume of data that needs to be transported and processed by the multi-process based O^2 system the data transport cost should be minimized. To achieve optimal performance the in-memory representation of data should allow a zero-copy approach. This is only possible when individual data units are POD type (plain old data), i.e. a scalar or a class type with no virtual base classes or functions that has no members of reference type or an array of such type. For full POD definition see e.g.[7].

As discussed, a zero-copy approach also introduces a tight relationship between the data and the transport layer, which should be handled in the interface design as transparently as possible.

3 O^2 data flow

In this section: overview of the main data flow and the requirements on the data layout and formats needed to achieve the flow:

3.1 First level processors

- CRU data in FLP memory.
- Subtimeframe building.
- FLP processing including quality control.

3.2 Event processing nodes

- Time frame building.
- reconstruction and calibration.
- compressed time frames.

3.3 Quality control (QC)

3.4 Asynchronous processing on the O^2 farm

- Event summary data.
- AOD format.
- Persistent data format(s).

3.5 Analysis facilities

- Data analysis requirements.

3.6 Requirement summary

4 Message queue based communication

The move to a multi-process, message queue based system is a large departure from the current monolithic ALICE offline data and computing model. It is, however, conceptually close to the modular processing model of the ALICE HLT. The decentralized messaging queue concept of FairMQ simplifies it by integrating the data transport layer into a library and eliminating the separation between dedicated transport and processing devices (or processes). In contrast to the managed and synchronized data flow in the HLT framework, message queue based systems are fully asynchronous. Messages are asynchronously queued on both the sending and receiving ends and multiplexed between the different connections without queue ordering guarantees. Devices receiving data are presented with an input queue where messages originating from different sources are interleaved in an non-deterministic order. For dealing with multiple logically associated data buffers, additional lower level abstractions are needed to maintain their association across transport boundaries without incurring excessive overhead.

Vectorized IO (also referred to as scatter/gather IO in the POSIX specification in [8]) is an important feature when dealing with multiple data buffers as it allows, in principle, to avoid the cost associated with serializing a data set into a single IO buffer before scheduling its transport. A form of vectored IO is provided in FairMQ as multi-part messages. A multi-part message consists of a structure (a vector) referencing multiple independent buffers. The multi-part message (and the referenced buffers) is transported atomically while preserving the ordering of the references.

The multi-part approach, in addition to minimizing the resource strain associated with IO buffer construction and/or data serialization, also by construction reduces the need for (re-)synchronisation and event building; data fragments once associated to a single logical unit (e.g. a time frame) remain that way throughout the entire chain regardless of the data transport topology. Another benefit is that additional data parts can be attached or removed by other processing devices without copy overhead at any point of the processing chain.

5 O² message structure

Each buffer associated to a single data set (e.g. a time frame) is described by metadata containing the information about e.g. the content type of the associated buffer, its origin and serialization strategy.

The O² message consists of a sequence of metadata-payload pairs of message parts contained within a multi-part message. In order to allow efficient payload classification, each buffer is described by metadata contained in a separate message part. Since FairMQ preserves the ordering of the parts, the natural choice is to precede each data part with the associated metadata part in the message, as illustrated in figure 1.

Since the O^2 message is a well-defined vector of references to the transported buffers, navigation is straightforward as it only involves access to the metadata via well-defined locations within the vector.

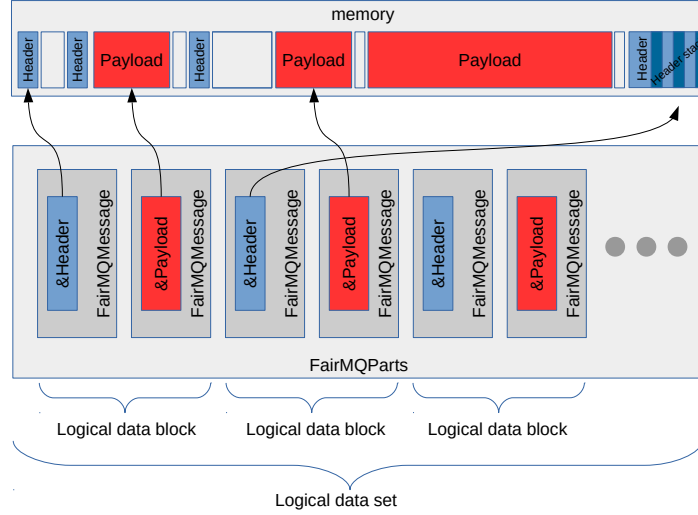


Figure 1: The O^2 message structure containing a data set. Logical data blocks consisting of payload and metadata parts are contained in a multi-part FairMQ message (bottom). The buffers associated to the message parts do not have to be contiguous in memory (top).

Storing the metadata and data payload in separate buffers offers several advantages:

- Since the metadata is separated from the payload already at transport level, efficient navigation is possible as only the (small) metadata parts need to be inspected.
- The size of the metadata buffer is not fixed enabling a scheme with flexible metadata content.
- The content of the data buffers, once produced by the hardware or a processing device, is immutable to other devices. Since the metadata is encapsulated in a separate buffer, it is possible to add additional information to the metadata (e.g. quality control status, contained in an additional header in the header stack, see section 6) with minimal cost and without modifying the payload downstream from the data producer.

6 Metadata format

The O^2 metadata consists of a contiguous buffer containing a sequence of headers (header stack). The byte representation of a header consists of a user defined body following a BaseHeader struct containing fields needed to:

- Verify that the following data (header body) belongs to an O^2 header.

- 208 • Define the size of the entire header.
- 209 • Flag whether another header follows this one in the stack.
- 210 • Verify the version of the header.
- 211 • Signal the type of the header.
- 212 • Signal the packing/serialization scheme of the metadata body carried
- 213 by the header.

214 The binary format of BaseHeader is fixed and can not be extended
 215 to guarantee consistent decoding of header information in the future. All
 216 metadata needs to follow the bytes representing the BaseHeader. The
 217 BaseHeader struct shall never be used directly, it only aids the construc-
 218 tion and decoding of metadata.

219 The header stack should contain at least the standard DataHeader
 220 struct describing the basic payload properties common to all payloads.
 221 The DataHeader representation starts with the BaseHeader followed by
 222 bytes representing:

- 223 • The functional data description uniquely determining the data type
- 224 contained in the payload.
- 225 • The payload serialization method (e.g. ROOT, FlatBuffers, none)
- 226 complementing the data description field.
- 227 • The origin of the data to identify the producer (e.g. detector system
- 228 or a software subsystem).
- 229 • A data type dependent 64 bit specification. This can be used by the
- 230 detectors to store e.g. the fine grained equipment ID like the link
- 231 number for raw data or cluster finder instance for clusters¹.
- 232 • The payload buffer size².

233 The binary format of DataHeader is fixed to ensure efficient access
 234 without the need for special decoding steps. The contents can only possi-
 235 bly be extended by appending additional data members and incrementing
 236 the version number to assure backward compatibility.

237 Other headers can be defined similarly to DataHeader and included
 238 in the header stack. Examples include trigger information headers for
 239 triggered detector data, object name headers for ROOT objects used in
 240 quality control etc.

241 6.1 Header memory layout

242 The memory layout of the BaseHeader and the derived DataHeader structs
 243 is compatible with the following definition:

```
244 struct BaseHeader
245 {
246     uint32_t    magicString;
247     uint32_t    headerSize;
248     uint32_t    flags;
249     uint32_t    headerVersion;
```

¹Based on HLT experience: most data types use some kind of fine grained ID. This field has been added here to avoid the overhead of a full header for what in most cases would be just one (64 bit) field.

²This is not strictly necessary online as the transport framework keeps track of the buffer sizes. Keeping this information in the header is useful for persistent storage and debugging purposes.

```

250     uint64_t    headerDescription;
251     uint64_t    headerSerialization;
252 };
253
254 struct DataHeader
255 {
256     BaseHeader baseHeader;
257     uint64_t    dataDescription[2];
258     uint32_t    dataOrigin;
259     uint32_t    reserved;
260     uint64_t    payloadSerializationMethod;
261     uint64_t    subSpecification;
262     uint64_t    payloadSize;
263 };

```

7 Data formats

The O^2 data model does not impose any limitations on the data types exchanged between devices. The only constraint from the data transport layer is that the transported payload must be contiguous in memory. In addition, the buffer for the payload should be allocated by the transport layer in a memory region appropriate for the chosen transport method. For higher level data types it usually means that they need to be serialized which penalizes performance (to a varying degree). A trade off has to be made between flexibility and performance based on the use case for a particular data type.

7.1 Data in memory

The recommendation is to use flat POD data types where performance and/or memory usage is critical, e.g. synchronous reconstruction and calibration, which needs to access and process all of the data. Outside of the critical path where relatively a low volume of data needs to be transported between devices, serialization schemes possibly impose acceptable overhead. The data model supports serialization schemes and facilitates the handling of serialization and deserialization of data transparently to user code.

7.2 Raw data formats

Raw data formats will be determined by the readout hardware of the detectors.

7.3 Reconstructed data formats

7.4 Persistent storage of time frame data

The in-memory data representation of header and payload buffers contains enough information to be stored on-disk directly as a sequence of buffers. POD data are, however, not suitable for long term storage and need to be transformed to a portable and extensible format.

In the modular O^2 design only a dedicated device would handle persistent storage making the translation steps from and to persistent storage

294 format transparent to other devices which only need to deal with the
295 in-memory format.

296 7.5 Quality control formats

297 The Quality control (QC) devices will produce ROOT based data.

298 7.6 AOD format

299 The contents of the AOD data used for end user physics data analysis is
300 being prepared by the physics community. Work is being done on develop-
301 ing the in-memory layout that will support efficient processing, pruning
302 and distribution in the message based O² environment. The struct-of-
303 arrays approach that is currently investigated also allows for transparent
304 extensions of the data content. The persistent storage of AOD still needs
305 to be investigated.

306 8 Interfaces

307 A minimal reference set of interfaces providing type-safe construction and
308 access to the metadata-payload protocol of the O² message is developed.
309 This complete, but minimal working set is meant to be augmented by
310 higher level interfaces based of the needs of the community.

311 8.1 Metadata access interface

312 Access to data members in the headers is protected by strongly typed
313 interfaces. Initialization can only be performed using predefined constants
314 unless the data member being set allows for arbitrary values explicitly;
315 consistency is then assured by the type system at compile time.

316 Inspection of header information contained in a metadata buffer is also
317 implemented in a type safe way.

318 8.1.1 Example

319 The user may check if a given buffer contains the desired header informa-
320 tion in a type safe way. If the desired header is of type DataHeader, a call
321 to:

```
322 DataHeader* header = Header::get<DataHeader>(buffer);
```

323 will yield a valid pointer if a header of type DataHeader is part of the
324 header stack inside the buffer pointed to by the buffer pointer.

325 8.2 Message navigation

326 The base functionality of a device running in the O² system is contained in
327 the FairMQDevice class. This is a generic class providing all the necessary
328 logic to support the O² system and data model, like control abstractions, a
329 state machine and multi-part messaging. FairMQ support for the multi-
330 part messaging is, however, unaware of the O² specific message layout
331 and the header to payload association. In order to enforce a consistent
332 handling of the metadata the necessary interface is implemented on top
333 of FairMQDevice as O2device.

Each device in the O² system should inherit from the O2device class in order to be able to use the data model. Interfaces are provided that insure consistency in the handling of the data and associated metadata.

8.3 Examples

To construct a message, the AddMessage(...) family of methods should be used to insure correct association of metadata to payloads, e.g. in the simplest case:

```
bool AddMessage( FairMQParts& message ,
                  AliceO2::Header::Block&& headerStack ,
                  FairMQMessagePtr payloadPart );
```

where in a single call the header stack to payload association is made and data is appended to the message.

After the message is constructed a call to the standard FairMQ send method suffices:

```
int64_t Send( const FairMQParts& parts ,
              const std::string& channel );
```

This queues the entire message for atomic delivery via a specified data channel.

For decoding the message on the receiving end functionality is provided to access the metadata and it's associated payload in a consistent way. The user implements a function that takes the metadata and the payload buffers as input, and uses that function via e.g. a call to:

```
template <typename T>
bool ForEach(
    O2message& parts ,
    bool (T::*memberFunction)(
        const byte* headerBuffer , size_t headerBufferSize ,
        const byte* dataBuffer , size_t dataBufferSize
    )
);
```

Inside the user provided function the metadata and the payload are decoded and processed.

References

- [1] Buncic P. and The ALICE Collaboration. *Technical Design Report for the Upgrade of the Online-Offline Computing System*. Tech. rep. ALICE-TDR-019 CERN-LHCC-2015-006. 2015.
- [2] *The Alice O2 software*. URL: <https://github.com/AliceO2Group>.
- [3] M. Al-Turany et al. "ALFA: The new ALICE-FAIR software framework". In: *Journal of Physics: Conference Series* 664.7 (2015), p. 072001. URL: <http://stacks.iop.org/1742-6596/664/i=7/a=072001>.

- 375 [4] P. Vande Vyvre A. Kluge. *The detector read-out in ALICE dur-*
376 *ing Run 3 and 4*. Tech. rep. DRAFT ALICE-TECH-2016-001.
377 URL: [http://svnweb.cern.ch/world/wsvn/alicedrrun3/](http://svnweb.cern.ch/world/wsvn/alicedrrun3/Notes/Run34SystemNote/detector-read-alice/ALICErun34_readout.pdf)
378 [Notes/Run34SystemNote/detector-read-alice/ALICErun34_](http://svnweb.cern.ch/world/wsvn/alicedrrun3/Notes/Run34SystemNote/detector-read-alice/ALICErun34_readout.pdf)
379 [readout.pdf](http://svnweb.cern.ch/world/wsvn/alicedrrun3/Notes/Run34SystemNote/detector-read-alice/ALICErun34_readout.pdf).
380 [5] URL: <https://root.cern.ch/>.
381 [6] URL: <https://google.github.io/flatbuffers/>.
382 [7] URL: <http://en.cppreference.com/w/cpp/concept/PODType>.
383 [8] URL: http://www.unix.org/version3/ieee_std.html.